

# CS61C – Machine Structures

## Lecture 5 –Memory Mangement

9/7/2007

John Wawrzynek

([www.cs.berkeley.edu/~johnw](http://www.cs.berkeley.edu/~johnw))

[www-inst.eecs.berkeley.edu/~cs61c/](http://www-inst.eecs.berkeley.edu/~cs61c/)

### Memory Allocation

---

- **Remember:**
  - Structure declaration does not allocate memory
  - Variable declaration does allocate memory
- **So far we have talked about several different ways to allocate memory for data:**
  1. Declaration at the beginning of a block  
`int i; struct Node list; char *string;`
  2. “Dynamic” allocation at runtime by calling allocation function (alloc).  
`ptr = (struct Node *) malloc(sizeof(struct Node));`
- **One more possibility exists:**

## Memory Allocation

---

```
int myGlobal;  
main() {  
  
}
```

- Data declared outside of any procedure (before main).
- Similar to declaration at beginning of a block, but has “global” scope.

## Where are these allocated?


---

- If declare **outside** a procedure, allocated in “static” storage
- If declare **inside** procedure, allocated on the “stack” and **freed when procedure returns**.
  - Note: `main()` is a procedure

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

## The Stack


- Stack frame includes:
  - Return “instruction” address
  - Parameters
  - Space for other local variables
- Stack frames form contiguous blocks of memory; stack pointer tells where top stack frame is
- When procedure ends, stack frame is tossed off the stack; frees memory for future stack frames



## Stack

- Last In, First Out (LIFO) data structure

```
main ()
{ a(0);
}
void a (int m)
{ b(1);
}
void b (int n)
{ c(2);
}
void c (int o)
{ d(3);
}
void d (int p)
{
}
```



## C Memory Management

---

- C has 3 pools of memory
  - **Static storage**: global variable storage, basically permanent, entire program run
  - **The Stack**: local variable storage, parameters, return address (location of "activation records" in Java or "stack frame" in C)
  - **The Heap** (dynamic storage): data lives until deallocated by programmer
- C requires knowing where objects are in memory, otherwise things don't work as expected
  - Java hides location of objects

CS 61C L05 C Memory Management (7)

Wawrzynek Spring 2007 © UCB

## The Heap (Dynamic memory)

---

- Large pool of memory, **not** allocated in contiguous order
  - back-to-back requests for heap memory could result in blocks very far apart
  - where Java **new** command allocates memory
- In C, specify number of **bytes** of memory explicitly to allocate item

```
struct int *iptr;  
iptr = (int *) malloc(8*sizeof(int));  
/* malloc returns type (void *),  
so need to cast to right type */
```

- **malloc()**: Allocates raw, uninitialized memory from heap


CS 61C L05 C Memory Management (8)

Wawrzynek Spring 2007 © UCB

## Typical C Memory Management

◦ A program's **address space** contains 4 regions:

- **stack**: local variables, grows downward
- **heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward
- **static data**: variables declared outside main, does not grow or shrink
- **code**: loaded when program starts, does not change



*For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory*


CS 61C L05 C Memory Management (9)

Wawrzynek Spring 2007 © UCB

## Intel 80x86 C Memory Management

◦ A C program's 80x86 **address space**:

- **heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward
- **static data**: variables declared outside main, does not grow or shrink
- **code**: loaded when program starts, does not change
- **stack**: local variables, grows downward



CS 61C L05 C Memory Management (10)

Wawrzynek Spring 2007 © UCB

## Memory Management

---

- How do we manage memory?
- **Code, Static storage are easy:**  
they never grow or shrink
- **Stack space is also easy:**  
stack frames are created and destroyed in last-in, first-out (LIFO) order
- **Managing the heap is tricky:**  
memory can be allocated / deallocated at any time

## Heap Management Requirements


---

- Want `malloc()` and `free()` to run quickly.
  - Want minimal memory overhead
  - Want to avoid ***fragmentation\**** –  
when most of our free memory is in many small chunks
    - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.
- \* This is technically called *external fragmentation*

## Heap Management

### ◦ An example


- Request R1 for 100 bytes
- Request R2 for 1 byte
- Memory from R1 is freed
- Request R3 for 50 bytes



## Heap Management

### ◦ An example

- Request R1 for 100 bytes
- Request R2 for 1 byte
- Memory from R1 is freed
- Request R3 for 50 bytes



## K&R Malloc/Free Implementation

---

- Look at Section 8.7 of K&R
  - Code in the book uses some C language features we haven't discussed and is written in a very terse style
- Each block of memory is preceded by a header that has two fields:  
**size** of the block and  
**a pointer to the next** block
- All **free blocks** are kept in a linked list, the pointer field is unused in an allocated block

## K&R Implementation

---

- `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- `free()` checks if the blocks adjacent to the freed block are also free
  - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block
  - Otherwise, the freed block is just added to the free list



## Choosing a block in `malloc()`

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
  - **best-fit**: choose the smallest block that is big enough for the request
  - **first-fit**: choose the first block we see that is big enough
  - **next-fit**: like first-fit but remember where we finished searching and resume searching from there

## Peer Instruction – Pros and Cons of fits

- A. The con of **first-fit** is that it results in many **small blocks** at the beginning of the free list
- B. The con of **next-fit** is it is **slower than first-fit**, since it takes longer in steady state to find a match
- C. The con of **best-fit** is that it **leaves lots of tiny blocks**

	ABC
1:	FFF
2:	FFT
3:	FTF
4:	FTT
5:	TFF
6:	TFT
7:	TTF
8:	TTT


## Tradeoffs of allocation policies

- **Best-fit:** Tries to limit fragmentation but at the cost of time (must examine all free blocks for each malloc). Leaves lots of small blocks (why?)
- **First-fit:** Quicker than best-fit (why?) but potentially more fragmentation. Tends to concentrate small blocks at the beginning of the free list (why?)
- **Next-fit:** Does not concentrate small blocks at front like first-fit, should be faster as a result.

## Example Dynamic Memory Allocator

### K&R malloc/free

- Each block begins with a header:
- Free list:
  - *Circularly linked*
  - Partially traversed by malloc and free
  - Block appear on list in increasing memory position
- *Next fit* algorithm for allocation
- New block taken from tail of next sufficiently large block
- Free merges blocks existing free block(s)



## K&R Allocator

```
typedef struct header {
    struct header * ptr; /* next free block */
    unsigned size;      /* size of this block */
} Header;
```

◦ **K&R uses a “union” type to force alignment. Complicates field extraction:  $p \rightarrow \text{size}$  becomes  $p \rightarrow s.\text{size}$**

• For simplicity here, we don't use the union.

◦ **Globals:**

```
static Header base;
                /* empty list to get started */
static Header *freep = NULL;
                /* start of free list */
```

CS 61C L05 C Memory Management (21)

Wawrzynek Spring 2007 © UCB

## K&R malloc

```
void *malloc(unsigned n bytes)
{
    Header *p, *prevp;
    Header *morecore(unsigned); /* used to get another large block from OS */
    unsigned nunits;

    nunits = (nbytes+sizeof(Header)-1)/sizeof(Header) +1;
    /* round up to allocate in units of sizeof(Header) */

    if ((prevp = freep) == NULL) { /* no free list yet */
        base.ptr = freep = prevp = &base; base.size = 0;
    }
    for (p = prevp->ptr; ; prevp = p, p = p->ptr) {
        if (p->size >= nunits) { /* big enough */
            if (p->size == nunits) /* exactly */
                prevp->ptr = p->ptr;
            else { /* allocate tail end */
                p->size -= nunits; p+= p->size; p->size = nunits;
            }
            freep = prevp; /* start next search here next time */
            return (void *) (p + 1); /* point past the header */
        }
    }
    if (p == freep) /* wrapped around free list */
        if ((p = morecore(nunits)) == NULL) return NULL; /* none left */
}
```

CS 61C L05 C Memory Management (22)

Wawrzynek Spring 2007 © UCB

## K&R free


```
void free(void *ap) /* keeps blocks ordered by address */
{
    Header *bp, *p;

    bp = (Header *)ap - 1; /* point to block header */


    for (p = freep; !(bp > p && bp < p->ptr); p = p->ptr)
        if (p >= p->ptr && (bp > p || bp < p->ptr)) break;
    /* freed block at start or end of arena */

    if (bp + bp->size == p->ptr) { /* join to next block */
        bp->size += p->ptr->size;
        bp->ptr = p->ptr->ptr;
    } else bp->ptr = p->ptr;
    if (p + p->size == bp) { /* join to previous block */
        p->size += bp->size;
        p->ptr = bp->ptr;
    } else
        p->ptr = bp;
    freep = p;
}
```

## “Core” memory? Some history.



Array of ferrite rings,  
“cores”, threaded with  
wires for writing and  
reading.



64 x 64 array of cores

*How much denser is modern semiconductor memory?*

## Automatic Memory Management

---

- Dynamically allocated memory is difficult to track – why not track it **automatically**?
- If we can keep track of what memory is in use, we can reclaim everything else.
  - Unreachable memory is called **garbage**, the process of reclaiming it is called **garbage collection**.
- So how do we track what is in use?
- Techniques depend heavily on the programming language and rely on help from the compiler

CS 61C L05 C Memory Management (25)

Wawrzynek Spring 2007 © UCB

## And in conclusion...

---

- C has 3 pools of memory
  - **Static storage**: global variable storage, basically permanent, entire program run
  - **The Stack**: local variable storage, parameters, return address
  - **The Heap** (dynamic storage): `malloc()` grabs space from here, `free()` returns it.
- `malloc()` handles free space with freelist. Three different ways to find free space when given a request:
  - **First fit** (find first one that's free)
  - **Next fit** (same as first, but remembers where left off)
  - **Best fit** (finds most “snug” free space)

CS 61C L05 C Memory Management (26)

Wawrzynek Spring 2007 © UCB

---

## Extras

## Slab Allocator

---

- **A different approach to memory management (used in GNU `libc`)**
- **Divide blocks in to “large” and “small” by picking an arbitrary threshold size. Blocks larger than this threshold are managed with a freelist (as before).**
- **For small blocks, allocate blocks in sizes that are powers of 2**
  - **e.g., if program wants to allocate 20 bytes, actually give it 32 bytes**

## Slab Allocator

---

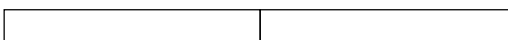
- Bookkeeping for small blocks is relatively easy: just use a *bitmap* for each range of blocks of the same size
- Allocating is easy and fast: compute the size of the block to allocate and find a free bit in the corresponding bitmap.
- Freeing is also easy and fast: figure out which slab the address belongs to and clear the corresponding bit.

## Slab Allocator

---

16 byte blocks: 

32 byte blocks: 

64 byte blocks: 

16 byte block bitmap: 11011000

32 byte block bitmap: 0111

64 byte block bitmap: 00

## Slab Allocator Tradeoffs

---

- Fast for small blocks.
- Slower for large blocks
  - But presumably the program will take more time to do something with a large block so the overhead is not as critical.
- Minimal space overhead
- No external fragmentation (as we defined it before) for small blocks, but still have wasted space!

## Internal vs. External Fragmentation

---

- With the slab allocator, difference between requested size and next power of 2 is wasted
  - e.g., if program wants to allocate 20 bytes and we give it a 32 byte block, 12 bytes are unused.
- We also refer to this as fragmentation, but call it *internal fragmentation* since the wasted space is actually within an allocated block.
- *External fragmentation*: wasted space between allocated blocks.



## Buddy System


---

- Yet another memory management technique (used in Linux kernel)
- Like GNU's "slab allocator", but only allocate blocks in sizes that are powers of 2 (internal fragmentation is possible)
- Keep separate free lists for each size
  - e.g., separate free lists for 16 byte, 32 byte, 64 byte blocks, etc.

## Buddy System

---

- If no free block of size  $n$  is available, find a block of size  $2n$  and split it into two blocks of size  $n$
- When a block of size  $n$  is freed, if its neighbor of size  $n$  is also free, combine the blocks in to a single block of size  $2n$ 
  - **Buddy** is block in other half of larger block



- Same speed advantages as slab allocator

## Allocation Schemes

---

◦ So which memory management scheme (K&R, slab, buddy) is best?

- There is no single best approach for every application.
- Different applications have different allocation / deallocation patterns.
- A scheme that works well for one application may work poorly for another application.